



SAL

SILICON AUSTRIA LABS

VLSI FOR MACHINE LEARNING

PEDRO JULIAN, DIEGO GIGENA, NICOLÁS
RODRÍGUEZ

CAE 2023, Córdoba



SAL

SILICON AUSTRIA LABS

5. VLSI IMPLEMENTATIONS

01

Basic Elements

02

Area and Power

03

Quantization

04

Advanced Elements

05

Architectures





SAL
SILICON AUSTRIA LABS



BASIC ELEMENTS

CMOS LOGIC


- ≡ Complementary Metal-Oxide Semiconductor
- ≡ MOS transistors used as switches
 - ≡ pMOS activated with LOW
 - ≡ nMOS activated with HIGH
- ≡ Logic gates composed from:
 - ≡ pMOS pull-up
 - ≡ nMOS pull-down




CMOS INVERTER





Cross section




Top view




CMOS INVERTER



Cross section




Top view




CMOS NAND/NOR





NAND Gate






NOR Gate



CMOS AND/OR






CMOS XOR/XNOR





CMOS XOR/XNOR

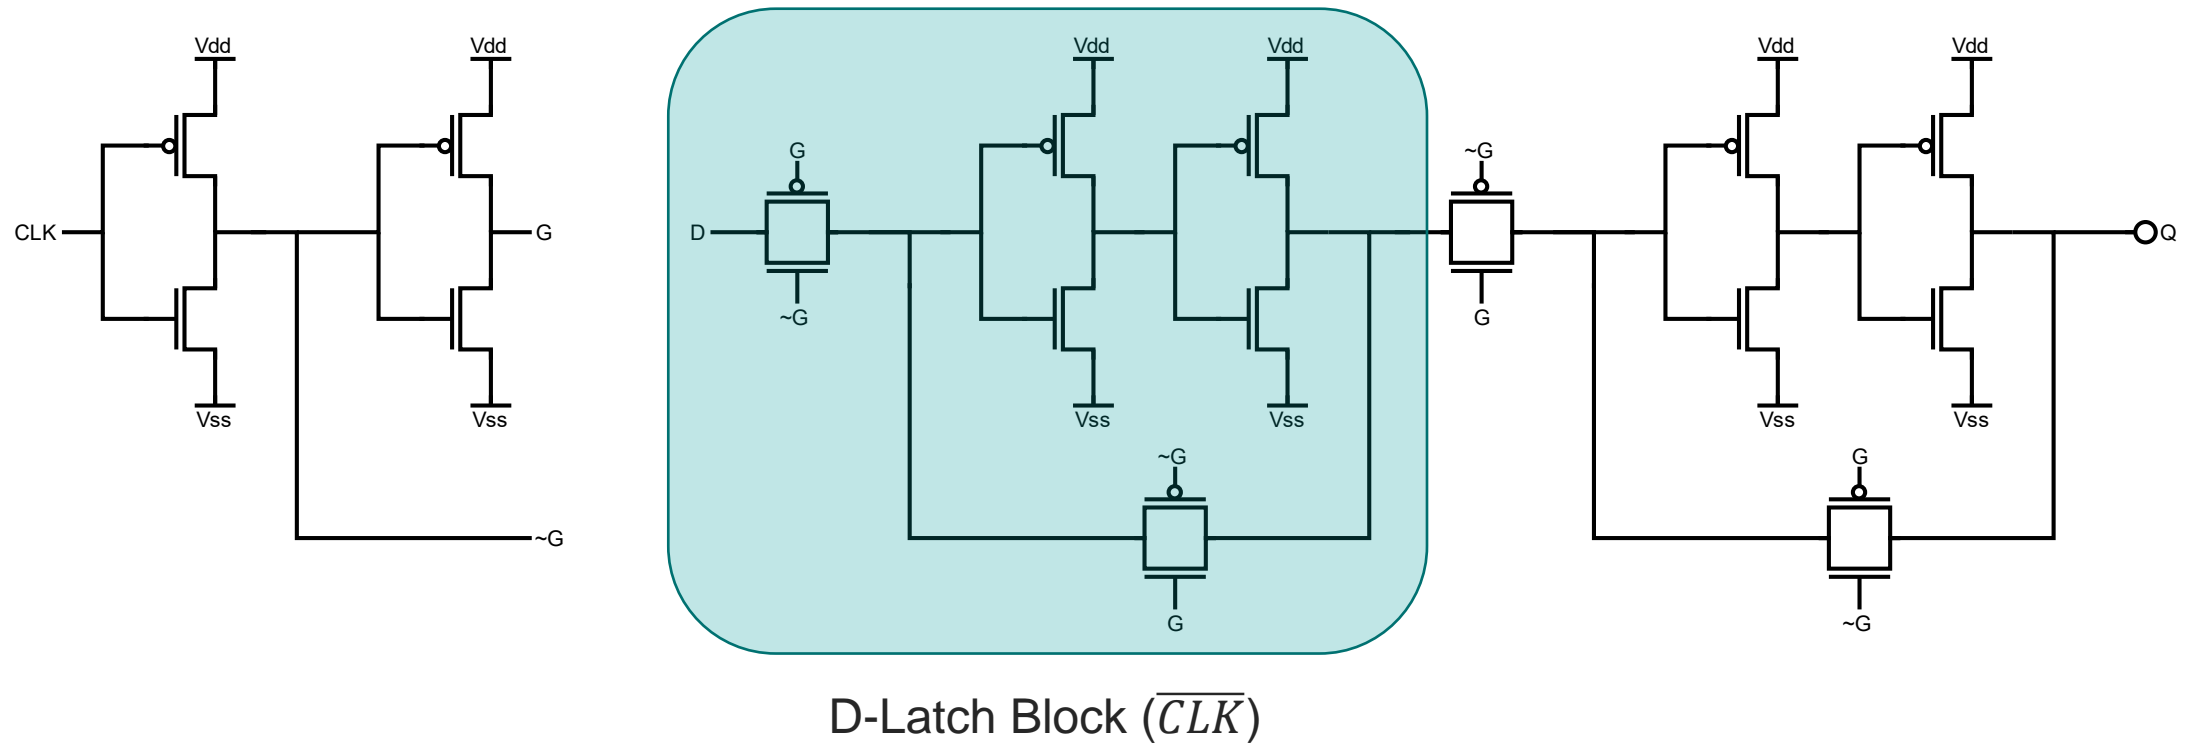
≡ XNOR has same number of transistors



CMOS D-LATCH





CMOS D-FLIP-FLOP




1BIT MULTIPLEXER

$$Y = A\bar{S} + BS$$




1-BIT ADDERS

Half Adder



Full Adder



SUMMARY BASIC ELEMENTS

≡ Number of transistors:

- ≡ INV: 2
- ≡ NAND/NOR: 4
- ≡ AND/OR: 6
- ≡ XOR/XNOR: 12
- ≡ D-Latch: 12
- ≡ D-Flip-Flop: 20+
- ≡ MUX: 20
- ≡ HA: 18
- ≡ FA: 42




SAL
SILICON AUSTRIA LABS

AREA AND POWER

AREA

$$Area = \sum Area_{cell} = \sum Pitch * Width_{cell}$$



AREA: EXAMPLE

≡ Compute area of Full Adder:

≡ XOR width: $2\ \mu m$

≡ AND width: $1\ \mu m$


≡ OR width: $1\ \mu m$

≡ Pitch: $2\ \mu m$

$$Area_{FA} = 2 * Area_{XOR} + 2 * Area_{AND} + Area_{OR}$$


$$Area_{FA} = 2\ \mu m * (2 * 2\ \mu m + 2 * 1\ \mu m + 1\ \mu m)$$

$$Area_{FA} = 14\ \mu m^2$$



LEAKAGE POWER

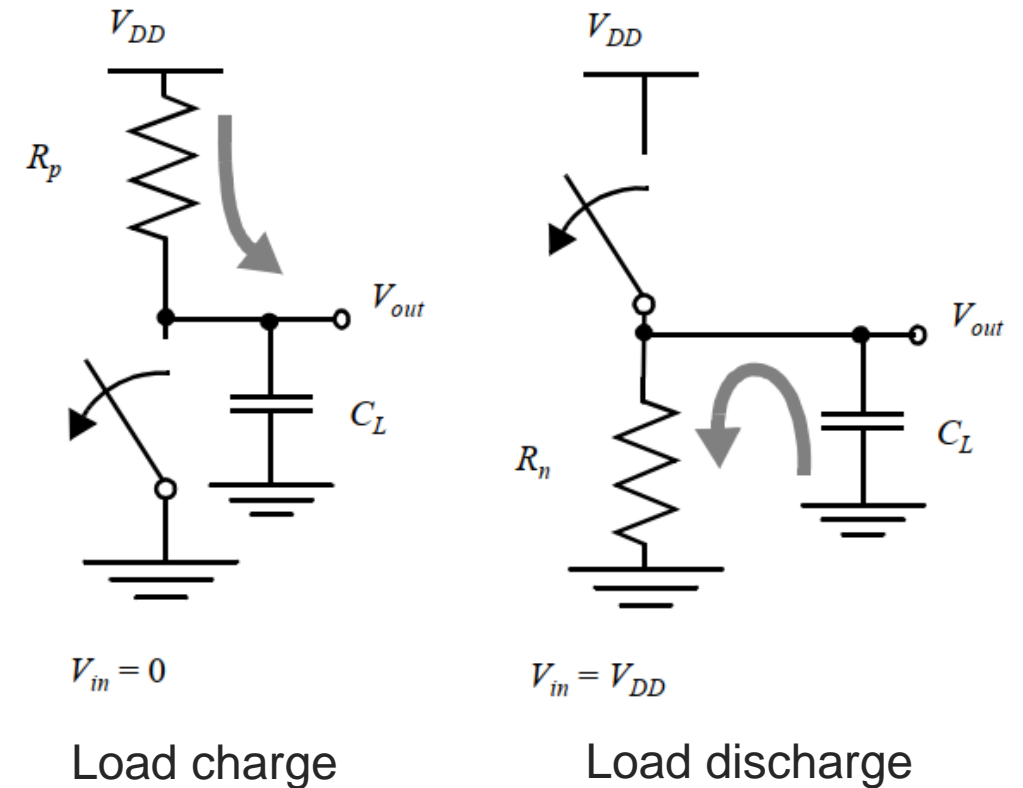
- ≡ Continuous power consumption
- ≡ Caused by:
 - ≡ Reverse bias current in parasitic diodes
 - ≡ Sub-threshold current
 - ≡ Gate oxide tunneling (thin gate oxide)
 - ≡ Gate-induced drain current



DYNAMIC POWER

- ≡ Power consumption during voltage transitions
- ≡ Depends on:
 - ≡ Transistors' sizes (capacitance)
 - ≡ Transition (rising or fall) times
 - ≡ Device threshold
 - ≡ Temperature
 - ≡ Switching activity ($\alpha_{0 \rightarrow 1}$ or $\alpha_{1 \rightarrow 0}$)
 - ≡ Node transitions per clock

$$Power_{dyn} = \frac{1}{2} \times \alpha \times C_L \times V_{DD}^2 \times f$$



POWER CONSUMPTION

≡ Simplest way of estimating dynamic power consumption:

$$Power = \sum Power_{cell} = \sum leak_{cell} + \left(f_{CLK} \times \sum avg_dyn_power_{cell} \times \alpha_{cell} \right)$$

POWER: EXAMPLE

≡ Compute Full Adder's power consumption:

≡ Typical Dynamic Power [$\mu W/MHz$]:

≡ XOR: 0.0040

≡ AND: 0.0022

≡ OR: 0.0025

≡ Leakage Power [μW]:

≡ XOR: 0.000016

≡ AND: 0.000014


≡ OR: 0.000013

≡ Clock frequency: 100 MHz

≡ Switching Activity:

≡ All gates have one transition every CLK

≡ $\alpha_{cell} = 1$



$$\begin{aligned} Power_{FA} &= 2 * (100 * 0.004 * 1 + 0.000016) \\ &\quad + 2 * (100 * 0.0022 * 1 + 0.000014) \\ &\quad + (100 * 0.0025 * 1 + 0.000013) \end{aligned}$$

$$Power_{FA} = 1.490073 \mu W$$

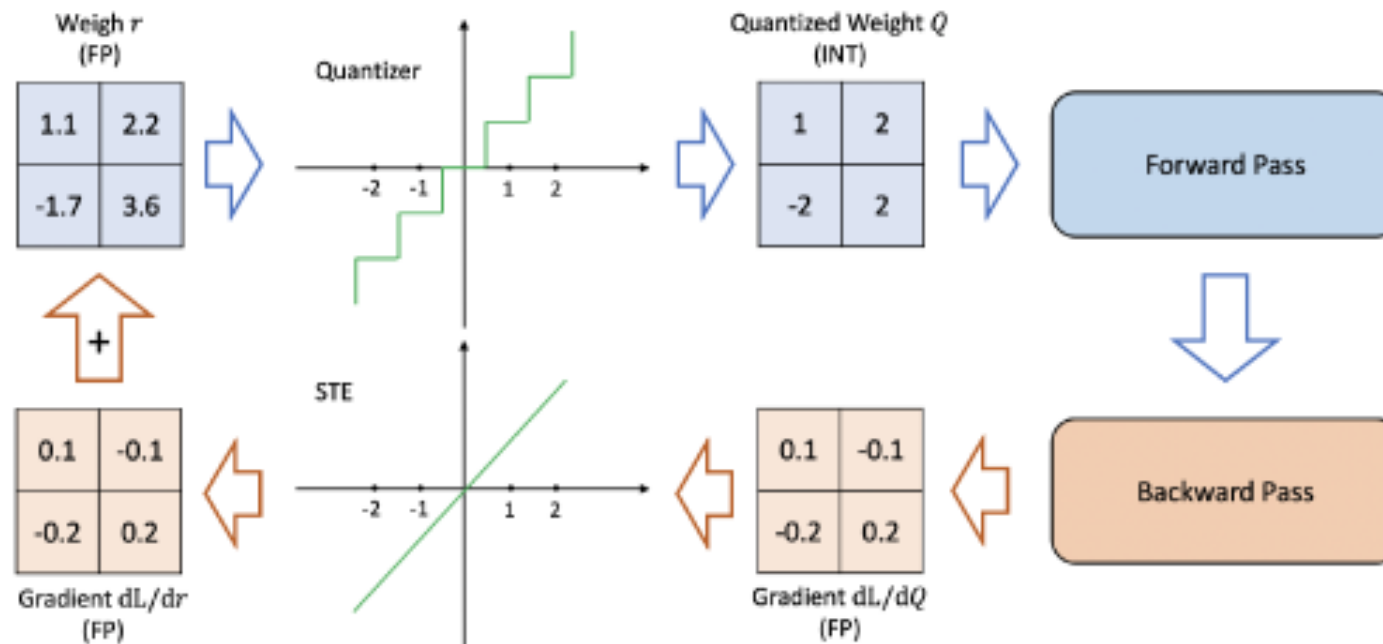


SAL
SILICON AUSTRIA LABS

QUANTIZATION


WHY QUANTIZATION

- ≡ Neural Networks run in CPU or GPU in general with FP operations
- ≡ VLSI/FPGA Accelerators mostly operate with INT values
- ≡ Quantization: mapping FP values into INT domain




FLOATING VS FIXED POINT

Floating Point 32b



$$Range_{fp32} = \begin{cases} [-3.40282347 \times 10^{38}, -1.17549435 \times 10^{-38}] \\ [0] \\ [1.17549435 \times 10^{-38}, 3.40282347 \times 10^{38}] \end{cases}$$




$$Range_{int8} = [-128, 127]$$

$$\left[\frac{-128}{Scale}, \frac{127}{Scale} \right]$$


TYPES OF QUANTIZATION

≡ Uniformity: Uniform vs Non-Uniform




TYPES OF QUANTIZATION

- ≡ Uniformity: Uniform vs Non-Uniform
- ≡ Symmetry: Symmetric vs Asymmetric (both Uniform)



Symmetric



Asymmetric

$$\text{Range} = [\alpha, \beta]$$

$$\text{Int}() = \begin{cases} \text{Round} \\ \text{Trunc} \\ \text{Floor} \end{cases} + \text{clip}$$

$$x_{int} = \text{Int}(S \times x_{fp})$$


$$S = \frac{2^{b-1}}{\max(|\alpha|, |\beta|)}$$

$$x_{int} = \text{Int}(S \times x_{fp}) - Z$$


$$S = \frac{2^{b-1}}{\beta - \alpha}$$

TYPES OF QUANTIZATION

-
- ≡ Uniformity: Uniform vs Non-Uniform
 - ≡ Symmetry: Symmetric vs Asymmetric (both Uniform)
 - ≡ Granularity: Tensorwise vs Channelwise



Tensorwise
(Layerwise)



Channelwise

TYPES OF QUANTIZATION

-
- ≡ Uniformity: Uniform vs Non-Uniform
 - ≡ Symmetry: Symmetric vs Asymmetric (both Uniform)
 - ≡ Granularity: Tensorwise vs Channelwise
 - ≡ Range Computation:
 - ≡ Dynamic (during runtime)
 - ≡ Static (after training)

QUANTIZED NEURON

≡ Example:

- ≡ Uniform
- ≡ Symmetric
- ≡ Tensorwise

$$y_{fp} = b_{fp} + \sum x_{fp} \times w_{fp}$$

$$\frac{y_{int}}{S_y} = \frac{b_{int}}{S_b} + \sum \frac{x_{int}}{S_x} \times \frac{w_{int}}{S_w}$$

$$S_y = S_b = S_x \times S_w$$

$$y_{int} = b_{int} + \sum x_{int} \times w_{int}$$

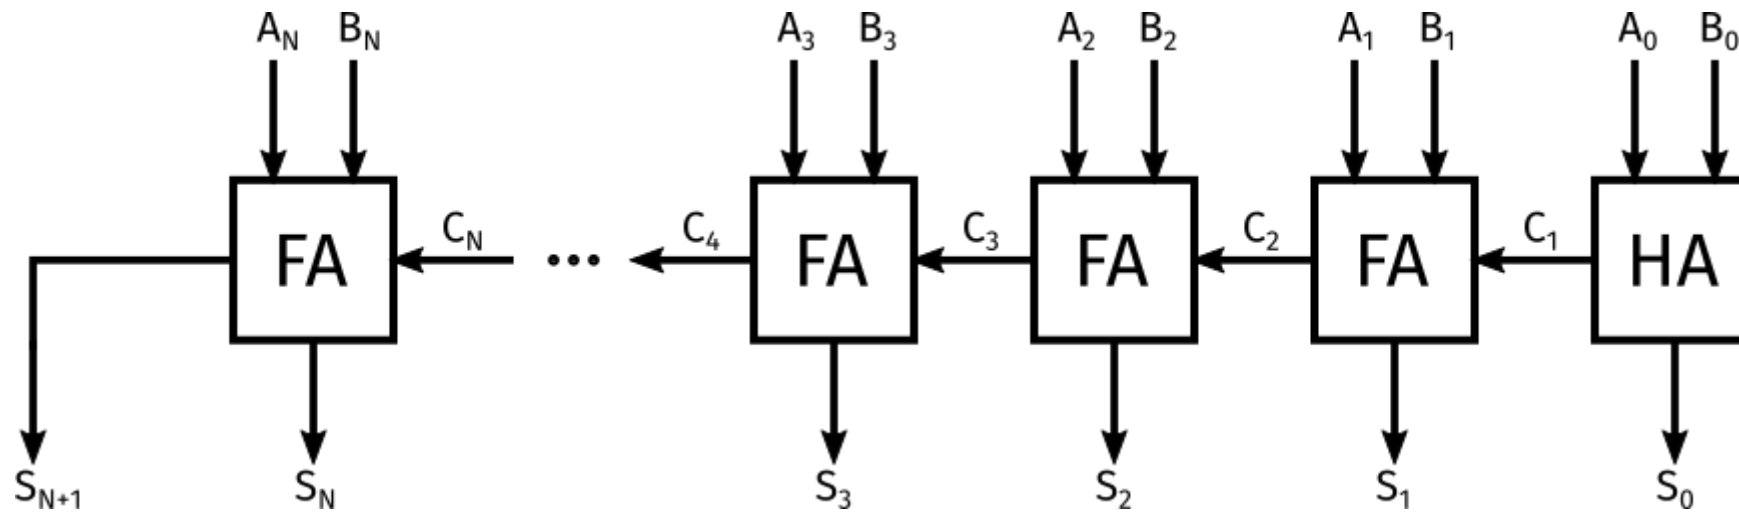
$$y_{fp} = \frac{y_{int}}{S_y}$$

$$x_{int}^{l+1} = y_{int}^l \frac{S_x^{l+1}}{S_y^l}$$

ADVANCED ELEMENTS

N-BIT ADDER

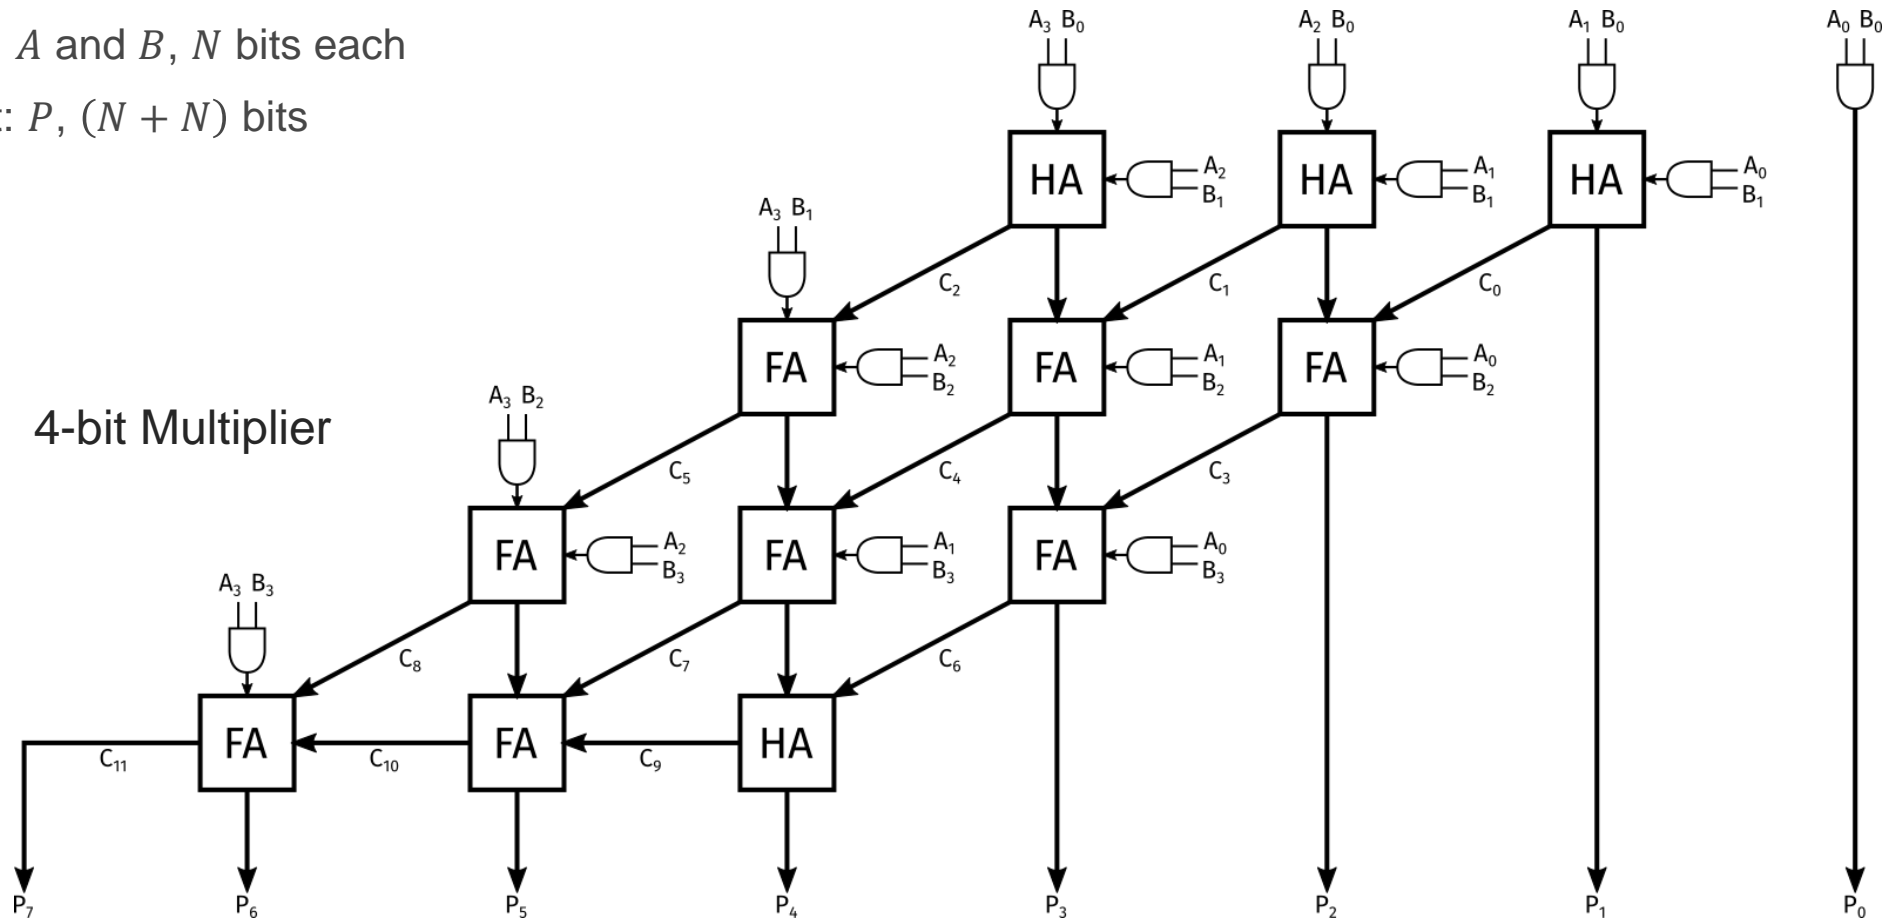
- ≡ Inputs: A and B , N bits each
- ≡ Output: S , $(N + 1)$ bits



N-BIT MULTIPLIER

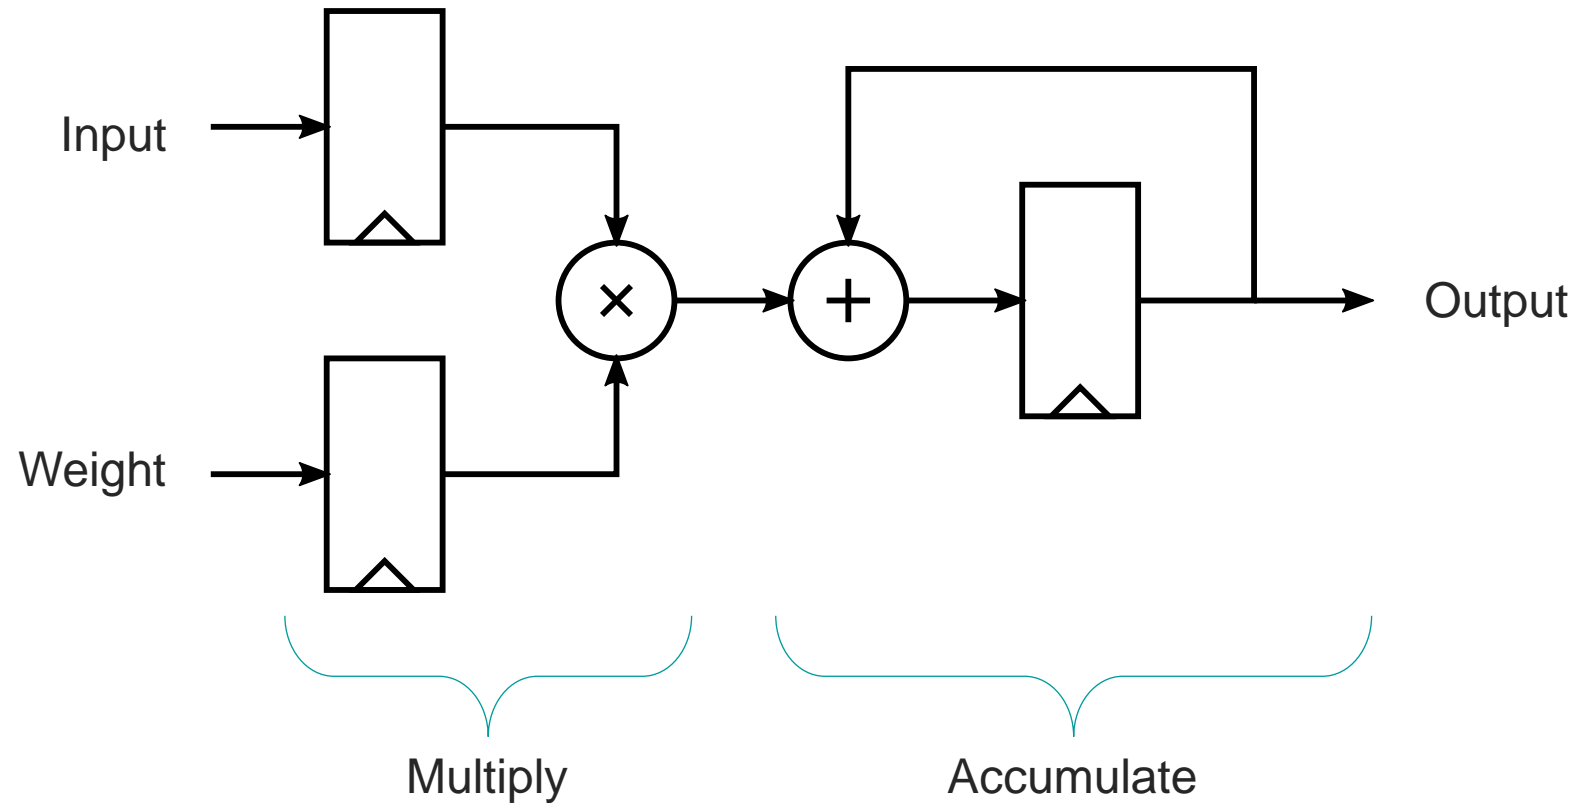
≡ Inputs: A and B , N bits each

≡ Output: P , $(N + N)$ bits




MAC

≡ Multiply and Accumulate operation



RELU OPERATOR

-
- ≡ Most Significant Bit (MSB) tells sign
 - ≡ Sign-magnitude
 - ≡ Ones' complement
 - ≡ Two's complement




COMPARATORS

≡ Equality ($A = B$)

≡ Inequality

≡ $A > B$

≡ $A < B$



1-bit Comparator

COMPARATORS

≡ Equality ($A = B$)

≡ Inequality

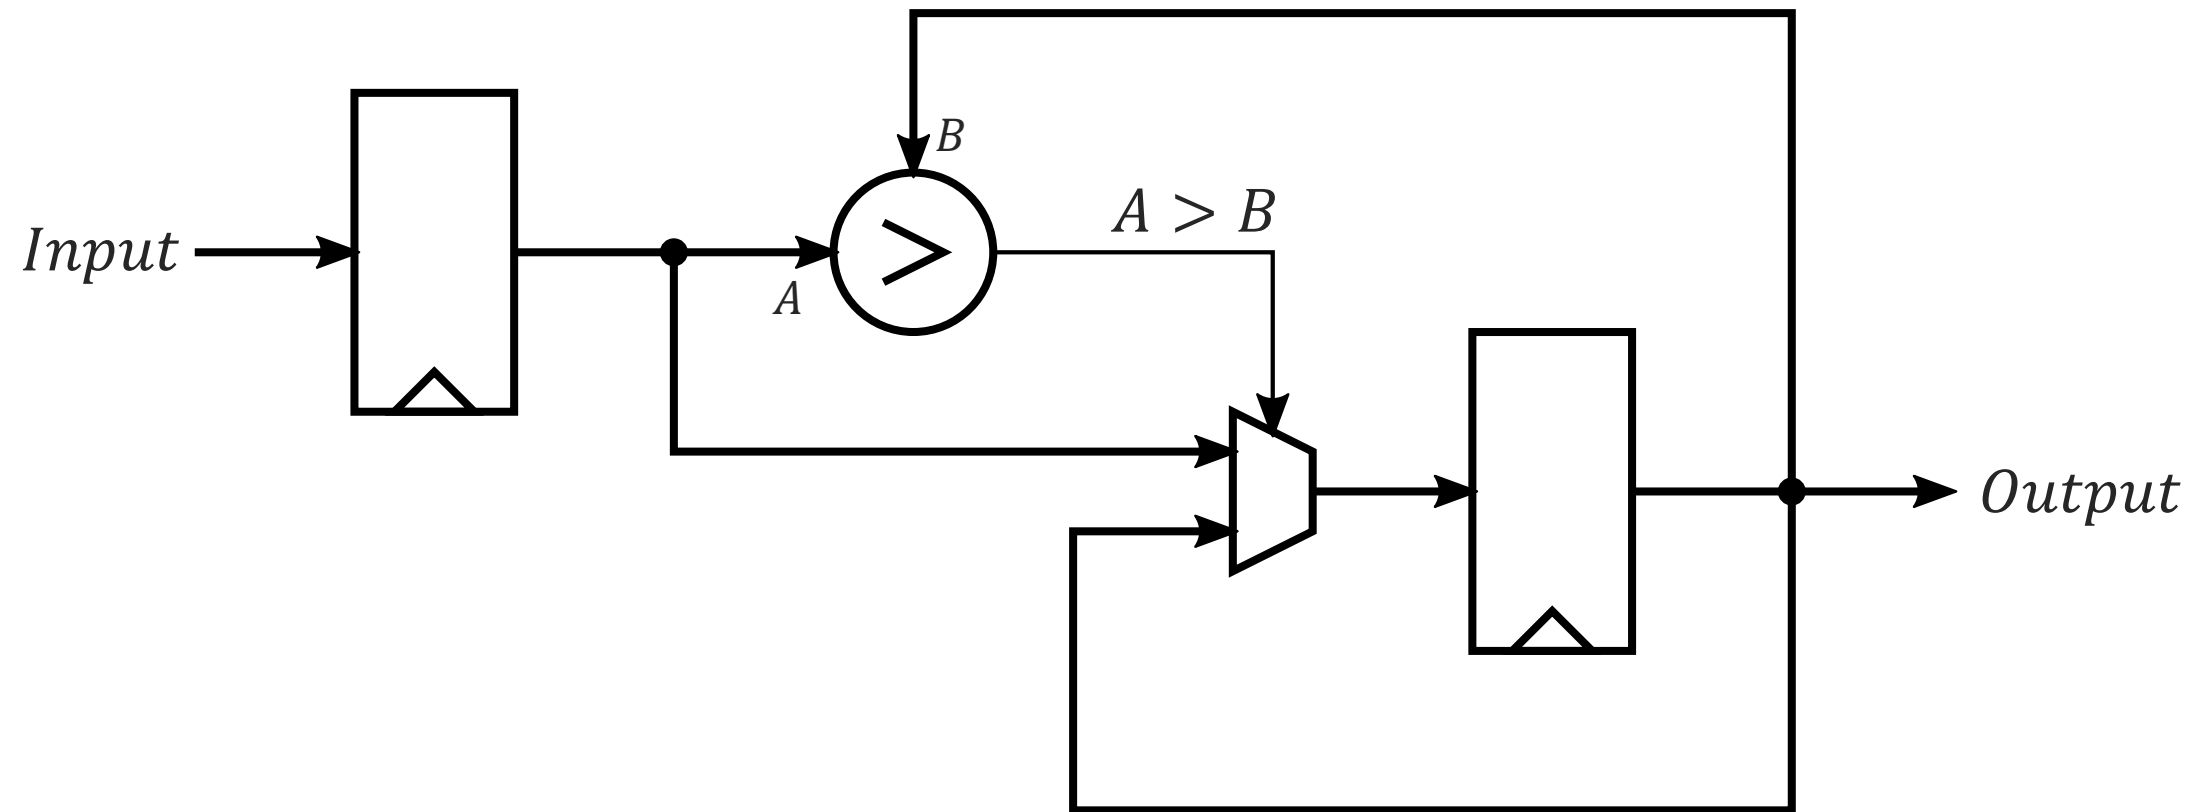
≡ $A > B$

≡ $A < B$

$$(A = B) = (A_N \oplus B_N) \cdots (A_1 \oplus B_1)(A_0 \oplus B_0)$$

$$x_i = (A_i = B_i) \left\{ \begin{array}{l} (A > B) = A_N \overline{B_N} + \cdots + (x_N \cdots x_2 A_1 \overline{B_1}) + (x_N \cdots x_2 x_1 A_0 \overline{B_0}) \\ (A < B) = \overline{A_N} B_N + \cdots + (x_N \cdots x_2 \overline{A_1} B_1) + (x_N \cdots x_2 x_1 \overline{A_0} B_0) \end{array} \right.$$

MAXPOOL OPERATOR





SAL
SILICON AUSTRIA LABS

ARCHITECTURES

PARALLEL COMPUTING


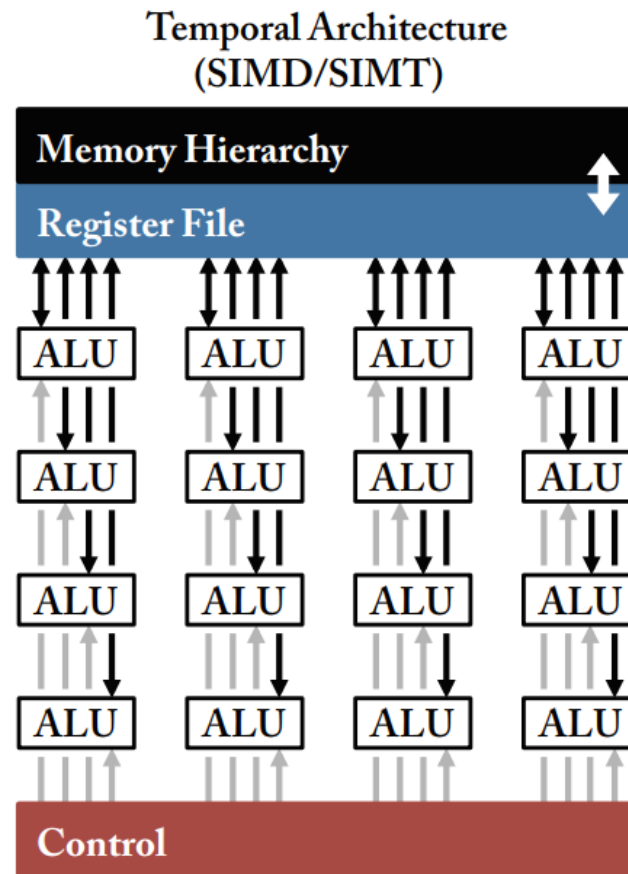
Two main paradigms

Temporal

- ≡ Centralized Control
- ≡ Centralized Memory
- ≡ Mostly CPUs/GPUs

Spatial

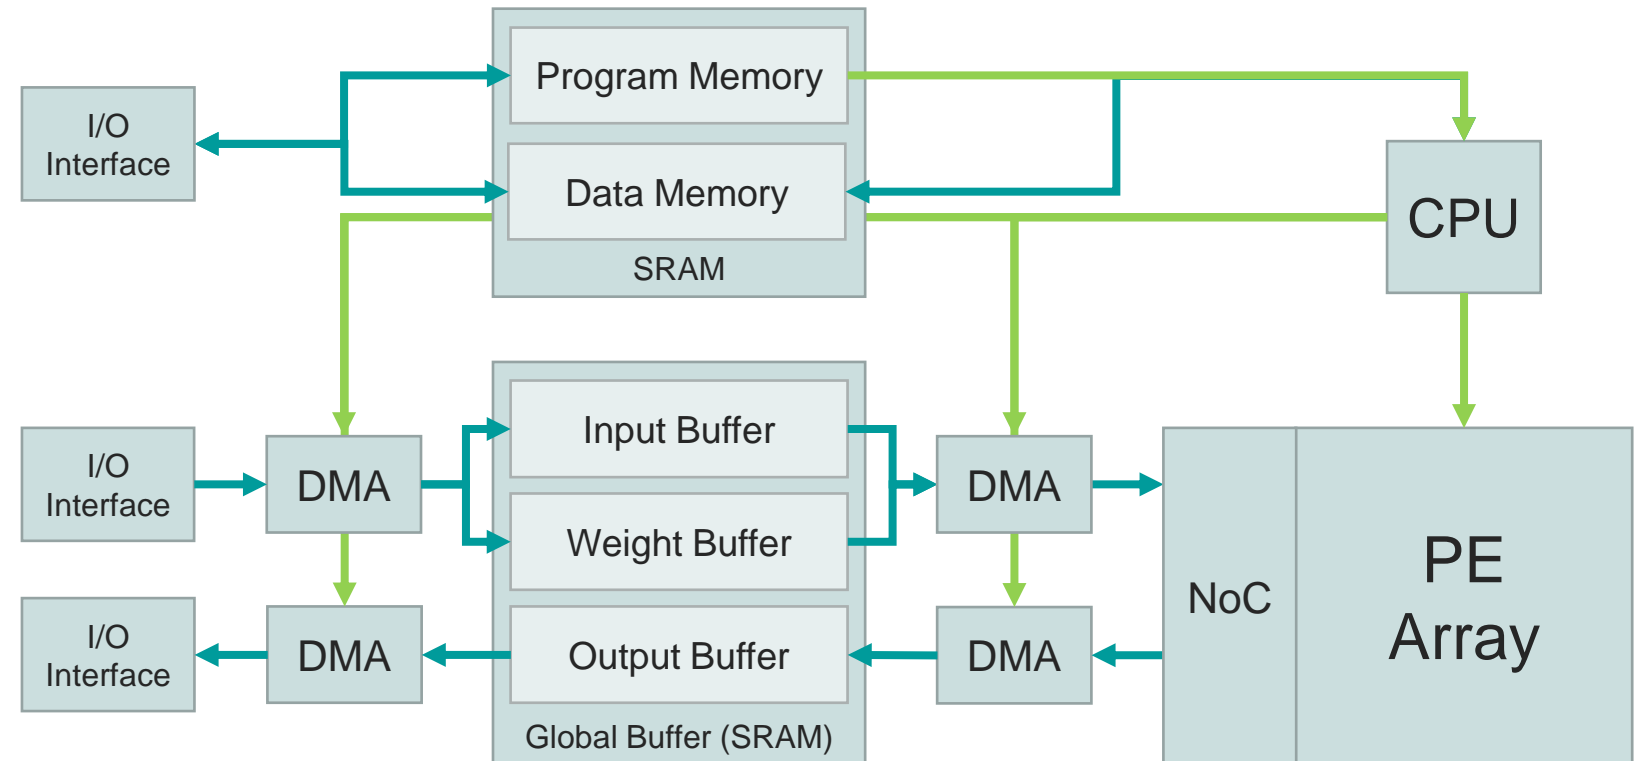
- ≡ Local Control
- ≡ Local Memory
- ≡ Inter-ALU Communication
- ≡ Mostly ASIC/FPGA
- ≡ ALUs + MEM + CTRL called PEs



ACCELERATOR BLOCKS

Accelerator Main Blocks:

- ≡ PE Array
- ≡ CPU (Global Control)
- ≡ SRAM
- ≡ Program
- ≡ Data
- ≡ Accelerator Buffers
- ≡ DMAs
- ≡ NoC
- ≡ Buffer – PE
- ≡ PE - PE
- ≡ I/O Interfaces



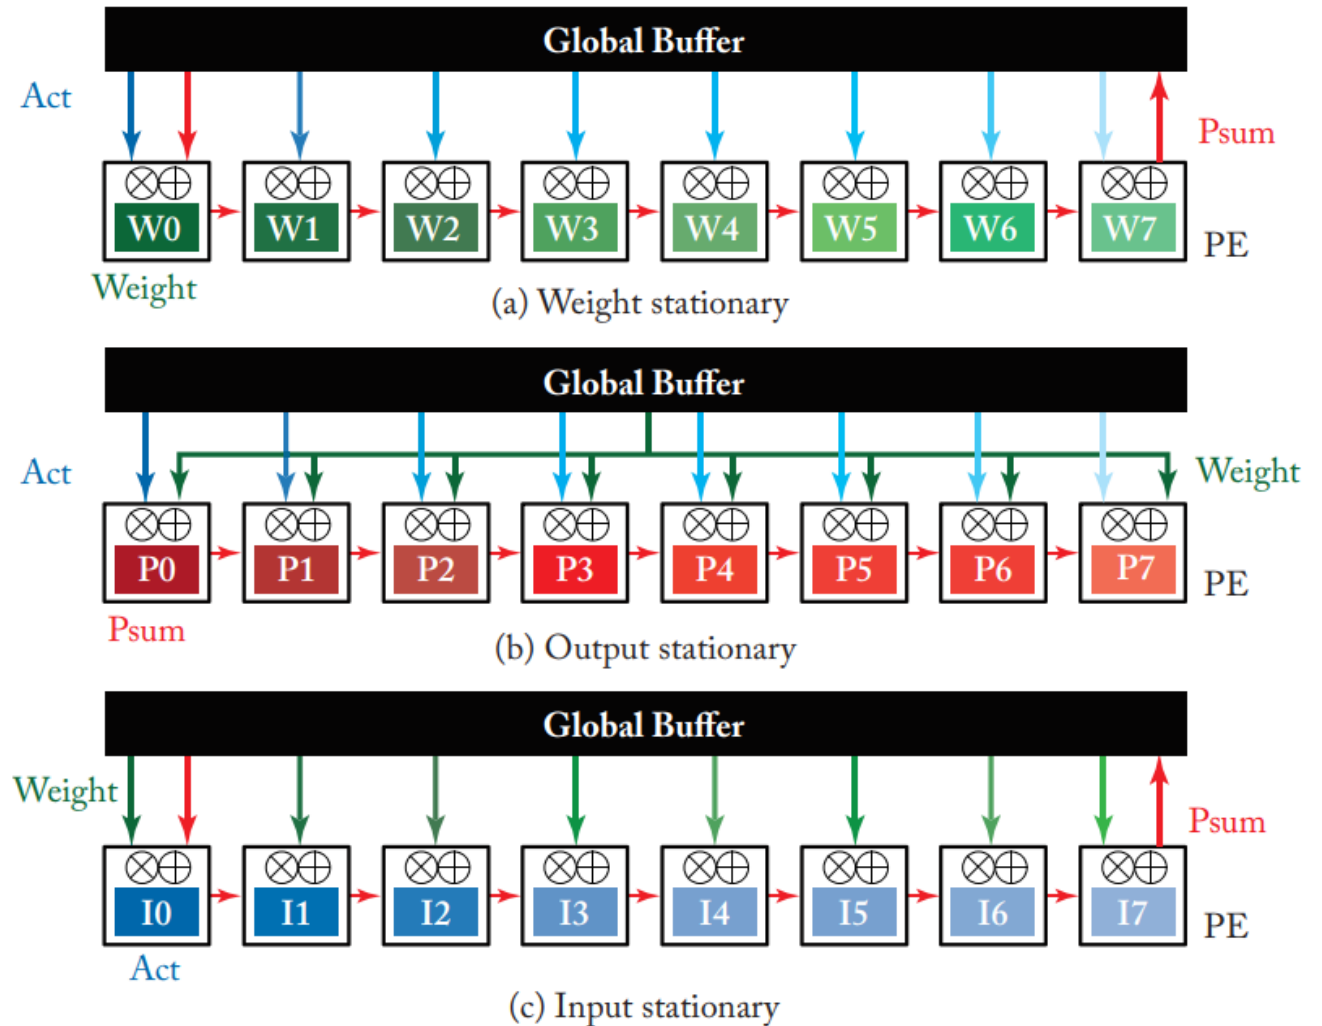
DATAFLOWS

- ≡ Exploit data reuse
- ≡ Dataflow styles:
 - ≡ Weight stationary
 - ≡ Input stationary
 - ≡ Output stationary


Conv Loop

```

for nn in OutH:
  for mm in OutW:
    for ll in OutCh:
      for kk in InCh:
        for jj in KerW:
          for ii in KerW:
            sum( $W \times In$ )
  
```




ACCELERATOR CHIPS



DagineuronV1

	DagineuronV1 (measured)	DagineuronV2 (simulated)
Technology	65 nm	65 nm
Power Supply	1.2 V	1.2 V
Clock freq.	100 MHz	100 MHz
Power active	2.37 mW	2.29 mW
TOPS/W	0.58	1.17



DagineuronV2



SAL
SILICON AUSTRIA LABS

QUESTIONS ?



SAL
SILICON AUSTRIA LABS

UNFOLD THE FUTURE